hw2-Classification

TAs
郭芳妤:clairekuo0217@gmail.com
楊依辰:sisia111062627@gapp.nthu.edu.tw
資電館743:若有問題請先來信預約

Hw2 Problem Discription

- Supervised multi-class classification problem (Credit Score)
- Given a data set
  - Training set with label
  - Testing set without label
- The dataset is transformed from a credit ranking dataset
  - 17 numeric features, 4 nominal features, 1 label
  - About 33643 cells become missing value
  - Our label is **CreditScore**
- Goal: predict the labels of testing data

Output Format

- Output your prediction to csv file with the following format and submit to kaggle
  Remember to output the first line

third_baseline.csv

	A	B
	Id	label
	0	2
	1	0
	2	1
	3	1
	4	0
	5	0
	6	0
	7	0

Id	label
0	2
1	0
2	0
3	0
4	0
5	0
6	0

Please remember you need only Id & label columns with 6762 rows (Id starts from 0) !

Evaluation

- We use **F1-score** = $2\times(precision\times recall)/(precision+recall)$
- **[update!!] We use macro-f1 in this homework**
- There are two leaderboards on Kaggle
  - Public: can be seen during competition
  - Private: can be seen after competition

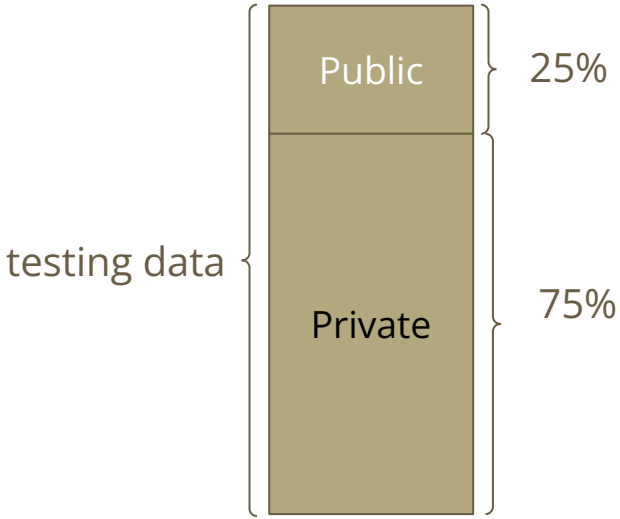

Hw2 Submission

- HW2 will be held on Kaggle
  - Please register a Kaggle account first
  - hw2 link: https://www.kaggle.com/competitions/nthu-2024datascience-hw02-classification
- Kaggle is a platform of
  - Machine learning competition
  - Sharing dataset
- Hw2 deadline: **2024.04.16** Tue. 23:59 (3 weeks)
- We will use the result on Kaggle to score this homework
  - **No need to hand in any files on eeclass**
  - Remember to **fill your Kaggle name** in the google form: https://forms.gle/h6Co4wwWp5GZwCPEA

Scoring and Rules

Scoring

- Use private leaderboard result for final scoring
- Baseline scores: We will score according to given 6 baseline scores

baseline	public	private
baseline-0	0.30155	0.33979
baseline-60	0.53815	0.53698
baseline-70	0.60884	0.60554
baseline-80	0.66737	0.68570
baseline-83	0.70888	0.73306
baseline-88	0.77935	0.80414

Scoring

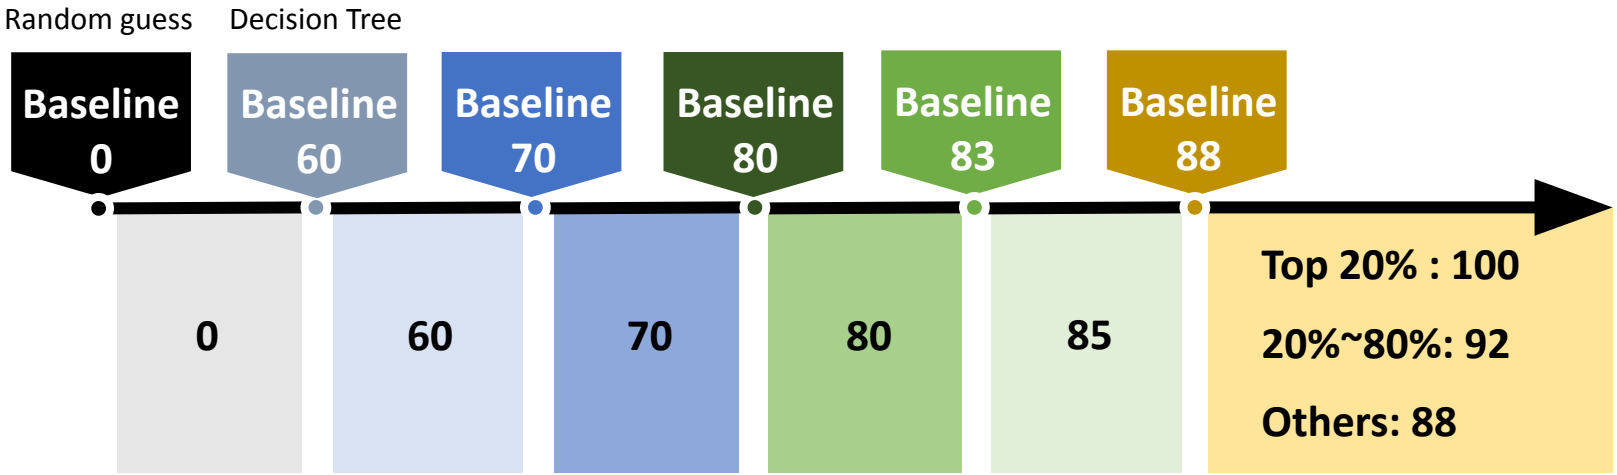

- You will get **0**, if your private score is between *baseline 0* and *baseline 60*
- You will get **60**, if your private score is between *baseline 60* and *baseline 70*
- You will get **70**, if your private score is between *baseline 70* and *baseline 80*
- And so on

Scoring

Baseline scores

- There are benchmarks on the leaderboard for reference

#	Team	Members	Score	Entries	Last
	baseline 88		0.77935		
	baseline 83		0.70888		
	baseline 80		0.66737		
	baseline 70		0.60884		
	baseline 60		0.53815		
	baseline 0		0.30155		

Rules

- You don’t have to submit the code!
- You can submit 20 times per day
- You can choose 4 predictions for final scoring
  - Kaggle will use the best one to be your final result
- No cheating!
- 若有問題，請寄信給負責助教。信件 title 請註明【DS-hw2-question】，並盡可能在信件內描述你遇到的困難，以方便我們協助你。
  - If you have any questions, please email the responsible teaching assistant. Please include "**[DS-hw2-question]**" in the email title and describe your difficulties as clearly as possible in the body of the email to facilitate our assistance.

How to submit and choose predictions

How to submit

- Click ***'Submit Predictions'*** button on the navigation bar

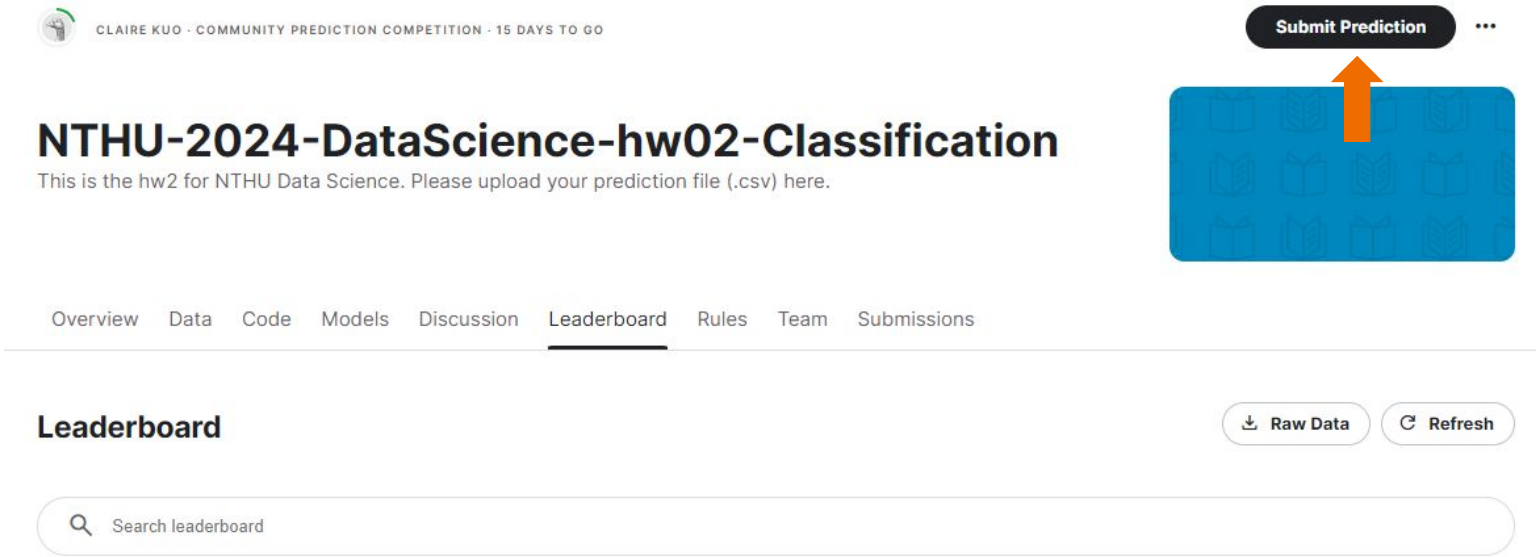

How to submit

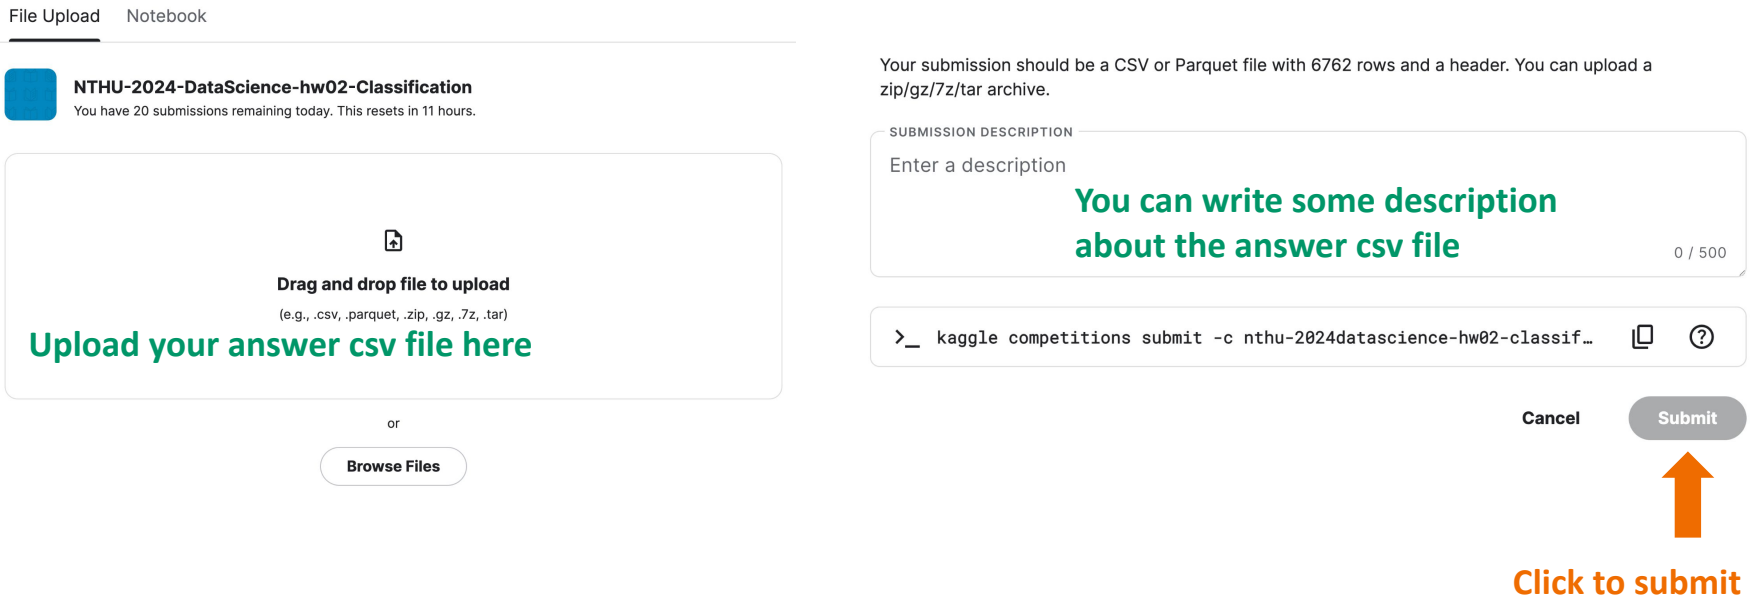

Hints

Hints

- How to read/write the file?

```
df = pd.read_csv("train.csv")
df_test = pd.read_csv("test.csv")
```

✓ 0.1s

```
y_pred = np.random.randint(3, size=len(df_test))
output = pd.DataFrame({'label': y_pred})

output.to_csv('myAns.csv', index_label='Id')
```

✓ 0.0s

Hints - Categorical features

`col_2`, `col_13`, `col_17`, `col_20` are categorical variables.
Please use the correct methods to handling those columns during training.

Hints

- Fillna with median in numeric features instead of 0

```
df[i] = df[i].fillna(median)
```

- Deal with data imbalance

```
from imblearn.over_sampling import SMOTE
sm = SMOTE(random_state=42)
X_train,y_train = m.fit_resample(X_train,y_train)
```

Hints

- Try different models
  - KNN, SVM, Logistic Regression, Random Forest ...

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
```

Finetune the model may achieve higher effect than the baseline 70 and 80

Hints

- **More techniques for better performance**
  - Feature selection (MI score)
  - Normalization
  - Dimension reduction (PCA, TSNE)
  - Try other different models
  - ...
- **We use private leaderboard as the final score**
  - Use public score to choose your model is dangerous
  - It’s better to perform validation

Packages you may use

- Scikit-learn
  - https://scikit-learn.org/stable/index.html
- Pandas
  - https://pandas.pydata.org/pandas-docs/stable/
- Imbalance learn (for over sampling and down sampling)
  - https://imbalanced-learn.org/stable/